PART I – Main Workshop Proposal *Pages 1–3, three-page limit*

AIDaR: Data Infrastructure and Benchmarks for Scientific AI

NeurIPS 2026 Workshop Proposal

Tagline (<140 chars): *From experimental data to reliable scientific AI systems.*

Abstract

Scientific AI is moving from curated prediction tasks toward systems that work directly with experimental data. Foundation models, retrieval systems, and agents are being asked to use experimental data, query scientific knowledge, run analyses, and support follow-up decisions. Scientific data must remain connected to samples, assays, protocols, measurements, workflow outputs, and the context needed to interpret results outside the lab or notebook where the data was produced.

AIDaR focuses on two problems in this transition. First, how should scientific data be organized so models and agents can use it: from raw measurements and analysis-ready files to relational or graph-structured knowledge, multimodal biomedical measurements, workflow systems, and large industrial datasets used for model training and wet-lab feedback? Second, how should evaluations determine where frontier models are reliable: across biological reasoning, practical data analysis, retrieval over structured scientific knowledge, assay-specific workflows, and deployed scientific AI systems?

The workshop will bring together researchers, industry scientists, and engineers building scientific data systems, foundation models, agents, biomedical evaluations, and industrial assay platforms. The program includes invited talks, contributed papers, demos, panels, and breakout groups focused on AI-facing data infrastructure, practical evaluations, and concrete lessons from scientific AI deployments.

1 Workshop Tracks

Scientific AI systems need data they can operate on and tests that show whether they used it correctly. AIDaR is organized around two tracks.

Track 1: AI-facing scientific data infrastructure. Scientific data must move from raw measurements into forms that models, retrieval systems, and agents can use. This includes linking files to samples and assay conditions, preserving the steps that produced an analysis-ready dataset, representing scientific relationships as tables, graphs, or knowledge networks, and exposing data through tools that scientists and agents can both use. Topics include data infrastructure for scientific foundation models, relational and graph-structured scientific data, multimodal biomedical measurements, assay-specific context, workflow systems, and large industrial datasets used for model training and wet-lab feedback.

Track 2: Benchmarks and evaluations for scientific AI. Scientific AI systems should be tested on tasks that resemble real scientific work. This includes choosing inputs, running analyses, checking controls, interpreting noisy results, and recovering biological or physical conclusions from data. Topics include biological reasoning and data-analysis tasks, retrieval and reasoning over structured scientific knowledge, evaluation of multimodal biomedical data, practical biology benchmarks, failure analysis, and studies of where model or agent performance changes across assays, platforms, workflows, or data organization choices.

Recent work on FAIR scientific data, dataset documentation, ML-ready metadata, AI-ready data, and scientific data-readiness pipelines motivates this workshop. AIDaR connects those efforts to scientific foundation models, retrieval systems, agents, and practical scientific AI evaluation [1, 3, 4, 5, 8, 9, 10, 6].

2 Why Now

Scientific AI is moving from models trained on curated datasets to systems that interact with experimental data, scientific literature, databases, and lab workflows. These systems need data that preserves sample, assay, protocol, measurement, and workflow context.

Modern scientific data systems are not organized for this style of use. Raw instrument data, sample tables, images, count matrices, workflow outputs, notebooks, and plots are often split across storage systems, LIMS/ELN tools, workflow runners, and lab-specific conventions. Large industrial datasets add further pressure: perturbation datasets, multiplexed screens, tissue-model systems, and wet-lab feedback loops must be connected to model-training datasets and analysis tools.

Evaluations now shape where frontier models can be used in scientific work. A model that performs well on general reasoning may still fail on a specific assay, platform, analysis workflow, or data modality. Practical evaluations help identify where agents can assist scientists, where human review is required, and where results are not reliable enough to support follow-up experiments or operational decisions.

AIDaR is timely because these problems sit between model development and scientific use. The workshop brings together researchers, industry scientists, and engineers building scientific data systems, foundation models, agents, biomedical evaluations, industrial assay platforms, and deployed analysis environments.

2.1 Target audience

AIDaR is for researchers, scientists, and engineers building AI systems that must work with real scientific data. The audience includes AI-for-science researchers, foundation-model and agent builders, scientific RAG researchers, computational biologists, biomedical ML researchers, benchmark and evaluation researchers, data-management and provenance researchers, and teams responsible for assay data, workflow systems, LIMS/ELN integrations, model-training datasets, and deployed analysis tools.

We expect a medium-sized workshop: **100–200 attendees** and approximately **40–50 submissions**. We will welcome both methods papers and practitioner reports, including scientific data-system case studies, workflow and tool demos, evaluation suites, failure analyses, assay metadata schemas, retrieval benchmarks, agent task suites, and lessons from deployed scientific AI systems.

2.2 Difference from related workshops

AIDaR sits between AI-for-science workshops, dataset venues, and reliable-ML workshops. The NeurIPS AI for Science series addresses broad methods and bottlenecks for AI-enabled discovery; ML4PS focuses on machine learning for physical sciences; FM4LS focuses on multimodal foundation models and LLMs for life-science data; Data-Centric ML and the NeurIPS Datasets & Benchmarks track focus on datasets, benchmarks, and data-centric ML methods; Reliable ML from Unreliable Data emphasizes missing, corrupted, shifted, adversarial, or strategic data; and ADSD 2025 at ICDM focuses on creating, curating, and benchmarking AI-ready scientific datasets [14, 15, 16, 17, 18, 19].

AIDaR focuses on what happens after scientific data is produced: how it is organized for models and agents, how it appears inside real analysis workflows, and how to tell whether an AI system used it correctly. Compared with broad AI-for-science workshops, AIDaR is more focused on data infrastructure. Compared with dataset venues, it is more focused on use rather than publication. Compared with reliable-ML workshops, it is more focused on scientific workflows, assays, and deployed analysis environments.

The recent AI-ready-data literature and ADSD/ICPP activity show growing interest in the topic; AIDaR would be the first NeurIPS workshop centered on AI-facing scientific data systems and practical evaluations for scientific AI [8, 9, 10].

2.3 Workshop History and Concurrent Submissions

This workshop has not previously been held at NeurIPS or another venue. No listed organizer is concurrently proposing another NeurIPS 2026 workshop.

3 Program Format, Speakers, and Diversity

3.1 Interactive one-day format

AIDaR is designed as a one-day in-person workshop within the 7–9 hour NeurIPS format. The program deliberately avoids a talks-only structure: invited talks are short, contributed work is central, posters/demos receive substantial time, and the afternoon breakout groups produce short written outputs for the post-workshop report. Any remote presentation will be used only for unforeseen emergencies and will remain within the NeurIPS one-hour cap.

Table 1: Draft one-day agenda. Confirmed speakers are listed in Section 3.2; remaining panel slots may be filled by additional confirmed panelists or accepted contributed work. All times are local and may shift by venue.

Time	**Session**	**Topic / AIDaR lens**
08:45–09:10	Arrival, question wall, opening remarks, COI/LLM policies, and workshop tracks	What experimental data must retain, and where evaluations show reliable AI use
09:10–10:00	Invited talks: Max Welling (University of Amsterdam), and Jure Leskovec (Stanford)	Scientific AI systems need AI-operable data: foundation models, materials discovery platforms, relational scientific data, and retrieval/reasoning over structured knowledge
10:00–10:30	Coffee + poster/demo previews	Early visibility into submitted benchmarks, tools, datasets, and case studies
10:30–11:20	Contributed spotlights: Scientific data systems and practical evaluations	Short talks on data schemas, experimental metadata, workflow interfaces, retrieval tasks, agent benchmarks, and failure cases from scientific AI systems
11:20–12:00	Panel 1: Benchmarking AI on practical biology. Harihara Muralidharan (LatchBio), Kexin Huang (Phylo/Biomni), Pablo Meher Rojas (IBM/Dream Challenge) and Jon Laurent (Edison Scientific)	How to construct biological and biomedical evaluations that approximate realistic scientific work, especially biological reasoning and practical data-analysis tasks.
12:00–13:15	Poster/demo session 1 + lunch	Community discussion around accepted papers, demos, tools, datasets, and emerging benchmark proposals
13:15–14:00	Invited talks: Marianna Rapsomaniki (University of Lausanne), and Chloe-Agathe Azencott (Mines Paris–PSL)	Biomedical data must preserve experimental meaning: assay design, sample and tissue context, treatment conditions, multimodal measurements, biological signal, and reproducibility
14:00–14:50	Contributed spotlights: analysis histories and deployment lessons	Short talks on workflow records, human review, failed analyses, and what scientific AI systems need to reproduce results
14:50–15:40	Poster/demo session 2 + coffee	Discussion of accepted tools, datasets, workflow systems, benchmark tasks, and lessons from deployed scientific AI systems
15:40–16:25	Breakout groups: AI-facing data infrastructure for large industrial datasets: ManifoldBio, Polyphron, Tahoe Therapeutics, and LigandAI **TO CONFIRM**	Small-group work on large perturbation datasets, multiplexed in vivo screens, tissue-foundry systems, training large models, and integrating wet-lab feedback
16:25–17:05	Panel 2: Engineering AI-operable assay data systems. Hoifung Poon (Microsoft Research) and additional industry panelists from TakaraBio, Vizgen, AtlasXOmics and Portal **TO DO**	Sample metadata, raw instrument data, analysis-ready files, workflow inputs and outputs, version history, and scientist-facing interfaces for deployed scientific AI systems
17:05–17:45	Breakout report-out, community roadmap, closing remarks, and post-workshop report plan	Next steps for AI-facing data systems, practical evaluations, and shared examples from scientific deployments

3.2 Speakers and Panelists

Confirmed invited speakers and panelists include Max Welling (University of Amsterdam), Marianna Rapsomaniki (University of Lausanne), Jure Leskovec (Stanford), Chloe-Agathe Azencott (Mines Paris–PSL), Additional panelists are being finalized; any unfilled panel slots will be replaced by accepted contributed work or moderated discussion among confirmed participants.

The invited-speaker program develops the workshop's two tracks. Max Welling and Jure Leskovec open with AI-operable scientific data systems: heterogeneous data for foundation models and discovery platforms, and relational/graph infrastructure for retrieval and reasoning. Marianna Rapsomaniki and Chloe-Agathe Azencott then ground the evaluation track in biomedical data: assay design, sample and tissue context, treatment conditions, multimodal measurements, biological signal, and reproducibility.

3.3 Panels

The two panels focus on the workshop's two tracks. Panel 1, "Benchmarking AI on practical biology," will focus on how to construct evaluations that approximate realistic scientific work, especially biological reasoning and data-analysis tasks that are useful in practice, rather than only ranking model accuracy on static benchmarks. Panelists include Harihara Muralidharan (LatchBio), Kexin Huang (Phylo/Biomni), Pablo Meher Rojas (IBM/Dream Challenge), and Jon Laurent (Edison Scientific).

Panel 2, "Engineering AI-operable assay data systems," will focus on the infrastructure required to make experimental data usable by deployed scientific AI systems: sample metadata, raw instrument data, analysis-ready files, workflow inputs and outputs, version history, and scientist-facing interfaces. Panelists include Hoifung Poon (Microsoft Research), with additional industry panelists from TakaraBio, Vizgen, AtlasXOmics, and Portal under discussion. Confirmed panelists are expected to attend in person.

3.4 Breakout groups

Breakout groups will focus on AI-facing data infrastructure for large industrial datasets, with expected participation from ManifoldBio, Polyphron, Tahoe Therapeutics, and LigandAI. Participants will discuss how to organize data from large perturbation datasets, multiplexed in vivo screens, and tissue-foundry systems so models and agents can use it. Topics include molecule and protein designs, perturbations, doses, donor and tissue context, sequencing or imaging readouts, assay outcomes, model-training datasets, and links from model predictions back to wet-lab follow-up. Groups will identify where current data organization breaks down as datasets move from experimental platforms into model training, retrieval, and agent-guided analysis, then report concrete recommendations for the post-workshop report.

3.5 Diversity, Inclusion, and Community Building

AIDaR is designed to bring together communities that rarely interact in the same venue: AI-for-science researchers, scientific data-infrastructure builders, benchmark and evaluation researchers, provenance and data-management experts, and practitioners deploying AI systems in high-stakes scientific settings. The organizing team and program committee span academia, industry research, biotechnology, scientific infrastructure, and data systems, with diversity across geography, seniority, discipline, institution type, and research perspective.

The workshop format is intentionally participatory rather than talks-only. Posters, demos, contributed spotlights, panels, and breakout groups are designed to give junior researchers, practitioners, tool builders, and domain scientists multiple opportunities to contribute. The CFP will explicitly welcome early-stage work, benchmarks, benchmark failures, lessons learned from unsuccessful AI data preparation efforts, tools, datasets, evaluation protocols, case studies, negative results, failed deployments, and postmortems of AI failures linked to data conditions.

We will also support community building beyond the workshop itself through lightweight mentoring for first-time authors, cross-domain collaboration, and breakout groups focused on readiness definitions, benchmark-design principles, and provenance-aware evaluation.

4 Logistics, review process, and compliance

4.1 Submissions, review, COI, and LLM policy

All submissions will be managed through OpenReview. We will accept non-archival short papers, extended abstracts, benchmarks/tools, demos, and position papers; the CFP will state the access policy and discourage work already accepted at the main NeurIPS conference or another archival ML venue. Invited speakers will be asked not to repeat main-conference or near-identical workshop talks.

Each submission will receive three reviews where possible. The PC will be sized so that no reviewer is assigned more than three papers; with 40–50 expected submissions, this requires 45–50 reviewers, with extra reviewers or area chairs if submissions exceed 70. Organizers will not give invited talks, submit papers, or shape decisions for conflicted submissions. Reviewers will be instructed not to use LLMs in reviewing, consistent with NeurIPS 2026 guidance.

4.2 Timeline

July 2026: Website and CFP launch.
August 2026: Submission deadline.
September 2026: Review, meta-review, and notifications.
October 2026: Final materials and public schedule.
December 2026: Workshop.

4.3 Location preference and technical needs:

Preferred order: **Paris, Atlanta, Sydney**; any venue is acceptable if needed. Technical needs are standard: projector/screen, podium microphone, Q&A microphone, panel microphones/table, poster boards, and standard recording/livestream support. No special hardware is required.

5 Expected Outcomes

Concrete deliverables of the workshop will include: (i) a public post-workshop report on open problems in AI-facing scientific data systems and practical evaluations; (ii) examples of data-system patterns for samples, assays, workflow runs, model-training datasets, and wet-lab feedback; (iii) recommendations for evaluating scientific AI systems in specific task contexts; (iv) a collection of failure cases from scientific AI deployments and benchmarks; and (v) a seed network of researchers, industry scientists, and engineers to develop future shared tasks and infrastructure examples.

PART II – Organizer Information *Pages 4–5, two-page limit*

6 Organizer Information

6.1 COI and Concurrent-Proposal Statement

Each listed organizer will confirm in writing, prior to OpenReview submission, that (a) they are not listed on more than two NeurIPS 2026 workshop proposals, (b) they understand that organizers cannot give invited talks at this workshop and will limit their on-stage role to brief introductions or panel moderation, and (c) they will not submit papers to this workshop, nor will their students, postdocs, or close collaborators. Organizers will not review or shape acceptance decisions about submissions from within their own organization (corporation-wide for large corporations).

6.2 Organizing Team

Michal Rosen-Zvi. *The Hebrew University of Jerusalem; IBM Research; michal.rosen-zvi@mail.huji.ac.il.* Director of AI for Healthcare and Life Sciences at IBM Research and Adjunct Professor of Computational Medicine at the Hebrew University. Her research focuses on machine learning for healthcare, scientific discovery, and biomedical AI. She leads multiple large-scale AI initiatives spanning pharmaceutical research, scientific reasoning, and AI-driven discovery systems.

Zoe Piran. *Stanford University; Roche / Genentech; piran.zoe@gene.com.* Postdoctoral Fellow at Genentech and Stanford University working at the intersection of machine learning and computational biology. Her research develops computational methods for understanding cellular dynamics from multimodal biological data and connects modern ML techniques with scientific discovery workflows.

Kenny Workman. *LatchBio; kenny@latch.bio.* Co-founder and CTO of LatchBio developing infrastructure and benchmarks for data analysis agents. His interests span engineering in biology, computing and AI.

Edaeni Hamid. *Roche / Genentech; hamid.edaeni@gene.com.* Machine learning and data-platform practitioner focused on deploying AI systems across complex scientific and enterprise data ecosystems. Her interests include AI data readiness, scientific data infrastructure, evaluation, and operationalizing AI systems in research environments.

Vladimir Ermakov. *Edison Scientific; vermakov@edisonscientific.com.* Researcher and entrepreneur building AI systems for scientific discovery, biomedical research, and foundation-model applications. His work focuses on scientific reasoning, evaluation, evidence synthesis, and AI systems that support hypothesis generation and decision-making in complex research environments.

Arindam Sett. *Roche / Genentech; sett.arindam@gene.com.* Principal Machine Learning Engineer at Genentech focused on AI data readiness, scientific AI platforms, retrieval-augmented systems, and agentic workflows. His work spans data foundations, knowledge systems, and AI infrastructure for scientific discovery, with a particular interest in understanding how data ecosystems influence downstream AI behavior.

6.3 Why This Team Can Deliver the Workshop

The team spans industry research labs and academia, with combined expertise across applied ML for biomedical and scientific data, foundation models for science, scientific data infrastructure, and trustworthy/data-centric AI. The team intentionally blends senior organizers with first-time workshop organizers (two members) to support junior involvement while maintaining organizational experience. Outreach is already underway, with anchor speakers confirmed across AI for Science, foundation models, scientific infrastructure, and biomedical AI. The organizing team has assembled an initial reviewer pool of 44 researchers spanning academia, biotech, scientific infrastructure, and industry research, providing immediate capacity for workshop review and community outreach. The workshop website, CFP, OpenReview venue, reviewer recruitment plan, and post-workshop artifact plan all have designated owners; OpenReview venue setup, CFP drafting, and reviewer recruitment have explicit owners on the team. The team is intentionally diverse across gender, institution type (industry research, academic, biotech), geography (US, Israel), and career stage.

6.4 Diversity of the Organizing Team

The organizing team is diverse across gender, institution type, geography, seniority, and discipline. The team spans pharmaceutical research, academic research, scientific infrastructure, biotech platforms, and independent scientific-AI organizations, with members across the US, and Israel. The organizers intentionally bridge scientific ML, biomedical AI, data systems, provenance, foundation models for science, and production AI infrastructure. The CFP and PC recruitment will explicitly extend beyond the organizers' immediate professional networks through university mailing lists, professional societies, social media, and cross-institutional outreach to broaden representation across geography, career stage, institution type, and research perspective.

6.5 Program Committee Plan and Roster

We are assembling a Program Committee of 40–50 reviewers before the CFP launch. Assuming 40–50 submissions and three reviews per submission, this target keeps the expected reviewer load at or below three papers per reviewer. If submissions exceed 50, we will expand the PC or appoint area chairs to maintain review quality and timely decisions.

The PC is organized around expertise clusters aligned with the workshop's two tracks: AI-facing scientific data infrastructure; evaluations for scientific AI; workflow systems and deployed analysis tools; biomedical and scientific AI applications; industrial assay platforms; reproducibility, failure analysis, and human review. Each submission will be assigned to reviewers with matching technical expertise, and COI checks will be performed through OpenReview before assignments are finalized. The named roster below reflects the current confirmed reviewer pool.

Current confirmed reviewer pool.

R01: J. Karin — HUJI
R02: K. Pang — Stanford
R03: M. Schaefer — Stanford
R04: K. Ponsin — Rockefeller
R05: A. Ghandehari — UC Irvine
R06: A. Gorla — UCLA
R07: A. Shafran — ETH Zurich
R08: S. Nair — Genentech
R09: A. Gupta — Stanford
R10: M. Bereket — Stanford
R11: A. Wu — Genentech
R12: N. Mudrik — Johns Hopkins
R13: R. Mintz — HUJI
R14: R. Friedman — HUJI
R15: W. Cheng — Carnegie Mellon
R16: Z. Zhu — UC Berkeley
R17: R. Lacombe — HCVC
R18: P. Ogden — Manifold Bio
R19: Z. Yang — LatchBio
R20: H. Bhasin — TwentyTwo Bio
R21: P. Pao-Huang — Stanford
R22: T. Kitak — ETH Zürich
R23: J. Liu — Genentech
R24: H. M. Khiem — University of Cincinnati
R25: M. Patel — Walmart
R26: O. Bazgir — Oracle
R27: J. Rossen — Harvard
R28: V. Shitov — Helmholtz
R29: N. Walter — CISPA
R30: A. Fomitchev — UC Santa Cruz
R31: S. Joshi — Intuit
R32: S. Kurant — Technion
R33: Y. Lee — Xaira
R34: Y. Shamay — Technion
R35: R. Littman — Genentech
R36: E. Peterfreund — Yale
R37: P. Gajare — UC Irvine
R38: N. Azran — HUJI
R39: M. Danziger — IBM
R40: M. Osman — Polyphron
R41: L. Nuzhna — Impetus Grants
R42: J. Wang — Harvard
R43: S. Alber — Stanford
R44: E. De Brouwer — Genentech

Additional reviewers will be recruited if needed.

6.6 Other Workshops the Organizers are Concurrently Proposing

No listed organizer is concurrently proposing another NeurIPS 2026 workshop.

References *Outside the 3+2 page limit*

References

[1] M. D. Wilkinson, M. Dumontier, I. J. Aalbersberg, G. Appleton, M. Axton, A. Baak, et al. The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data*, 3:160018, 2016. doi:10.1038/sdata.2016.18.

[2] Alfredo Gonzalez-Espinoza AI-ready Research Data 2026. https://guides.library.cmu.edu/ResearchDataForAI

[3] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi:10.1145/3458723.

[4] M. Pushkarna, A. Zaldivar, and O. Kjartansson. Data Cards: Purposeful and Transparent Dataset Documentation for Responsible AI. In *ACM FAccT*, 2022. doi:10.1145/3531146.3533231.

[5] M. Akhtar, O. Benjelloun, C. Conforti, L. Foschini, J. Giner-Miguelez, P. Gijsbers, S. Goswami, N. Jain, M. Karamousadakis, M. Kuchnik, et al. Croissant: A Metadata Format for ML-Ready Datasets. In *NeurIPS Datasets and Benchmarks Track*, 2024. https://arxiv.org/abs/2403.19546.

[6] K. Hiniduma, S. Byna, and J. L. Bez. Data Readiness for AI: A 360-Degree Survey. arXiv preprint arXiv:2404.05779, 2024. https://arxiv.org/abs/2404.05779.

[7] K. Hiniduma, S. Byna, J. L. Bez, and R. Madduri. AI Data Readiness Inspector (AIDRIN) for Quantitative Assessment of Data Readiness for AI. In *Proc. 36th Intl. Conf. on Scientific and Statistical Database Management (SSDBM)*, 2024. doi:10.1145/3676288.3676296.

[8] W. Brewer, P. Widener, V. Anantharaj, F. Wang, T. Beck, A. Shankar, and S. Oral. Data Readiness for Scientific AI at Scale. In *Proc. 54th Intl. Conf. on Parallel Processing Companion (ICPP Companion)*, 2025. https://arxiv.org/abs/2507.23018.

[9] W. Brewer, P. Widener, V. Anantharaj, F. Wang, T. Beck, A. Shankar, and S. Oral. Data readiness pipeline patterns for scientific AI at scale: Insights from climate, fusion, life sciences, and materials. *AI Magazine*, 2026. doi:10.1002/aaai.70056.

[10] N. Majithia, T. Carey-Wilson, E. Simperl, and N. Shadbolt. An actionable framework for AI-ready data. *AI Magazine*, 2026. doi:10.1002/aaai.70054.

[11] E. Bhardwaj, H. Gujral, S. Wu, C. Zogheib, T. Maharaj, and C. Becker. The State of Data Curation at NeurIPS: An Assessment of Dataset Development Practices in the Datasets and Benchmarks Track. In *NeurIPS Datasets and Benchmarks Track*, 2024. https://arxiv.org/abs/2410.22473.

[12] Y. Wu, L. Ajmani, S. Longpre, and H. Li. A Systematic Review of NeurIPS Dataset Management Practices. In *NeurIPS Datasets and Benchmarks Track*, 2024. https://arxiv.org/abs/2411.00266.

[13] Harvard Business Review Analytic Services. *Taming the Complexity of AI Data Readiness*. HBR Analytic Services Pulse Survey, sponsored by Cloudera, 2026. https://hbr.org/resources/pdfs/comm/cloudera/TamingtheComplexityofAIDataReadiness.pdf.

[14] AI for Science Community. The Reach and Limits of AI for Scientific Discovery: NeurIPS 2025 AI for Science Workshop. https://ai4sciencecommunity.github.io/neurips25.

[15] Machine Learning and the Physical Sciences Workshop Organizers. Machine Learning and the Physical Sciences, NeurIPS 2025. https://ml4physicalsciences.github.io/2025/.

[16] FM4LS Workshop Organizers. NeurIPS 2025 Workshop on Multi-modal Foundation Models and Large Language Models for Life Sciences. https://nips2025fm4ls.github.io/.

[17] Data-centric Machine Learning Research Community. Data-centric Machine Learning Research Workshop. https://dmlr.ai/.

[18] Reliable ML Workshop Organizers. Reliable ML from Unreliable Data: NeurIPS 2025 Workshop. https://reliablemlworkshop.github.io/.

[19] ADSD Workshop Organizers. The 1st Workshop on AI-Ready Data for Science Discovery at ICDM 2025. https://cnicds.github.io/ICDM2025/.

[20] S. Afzal, C. Rajmohan, M. Kesarwani, S. Mehta, and H. Patel. Data Readiness Report. In *IEEE Intl. Conf. on Smart Data Services*, 2020. https://arxiv.org/abs/2010.07213.